



Introduction to Assembly Language

ECE 3140/CS 3420 — EMBEDDED SYSTEMS

© 2016 José F. Martínez. Unauthorized distribution prohibited.

What is assembly language?


- **Assembly code:** Human-readable, quasi-isomorphic translation of machine code
 - Ok, but what is machine code?
- **Machine code:** Binary-encoded instructions describing a program
 - Directly executable by the processor



Machine code example (made up)

Processor:

- Fetches next instruction from program
- Decodes instruction
- Executes according to instruction
- Rinse and repeat



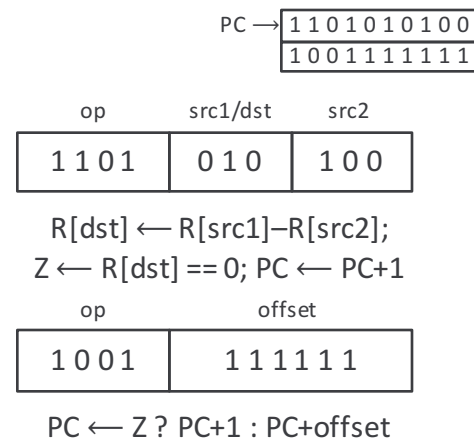
Machine code example (made up)

Assume:

- 16 different instructions
- 8 registers to store data (also PC, Z)
- Destination register is also source operand

Processor:

- Fetches next instruction from memory
- Decodes instruction
- Executes according to instruction
- Rinse and repeat



Update register with subtraction: used later for comparison

dependency of src1 with dst

dst is always updated, hence so is src 1 on the future

Assembly equivalent (made up)

Assume:

- 16 different instructions
- 8 registers to store data (also PC, Z)
- Destination register is also source operand

next:

SUB R2,R4
BNZ next

op	src1/dst	src2
SUB	2	4

Processor:

- Fetches next instruction from memory
- Decodes instruction
- Executes according to instruction
- Rinse and repeat

$R2 \leftarrow R2 - R4; Z \leftarrow R2 == 0; PC \leftarrow PC + 1$


op	offset
BNZ	-1

$PC \leftarrow Z ? PC + 1 : PC - 1$



Who programs in assembly?

- Nowadays, mostly people that enjoy pain and suffering
 - Ok, some low-level tasks best in assembly
- ECE 3140/CS 3420 students (for a few weeks at least)
- Compilers/interpreters *extremely good* at generating fast machine code from high-level languages
 - Tendency for bloated executables (e.g., libraries) →
 - Not always fastest (e.g., critical code block)



iDownload Blog



So why study assembly?

- Understand **hardware-software interface**
 - What functionality does the hardware provide?
 - How are high-level language constructs supported?
 - Subroutines, recursion
 - How are system services provided?
 - Dynamic allocation of variables
 - Interaction with I/O devices
 - Multitasking
- Build “bare-metal” (embedded) systems
 - Minimize code bloat; speed up critical code blocks



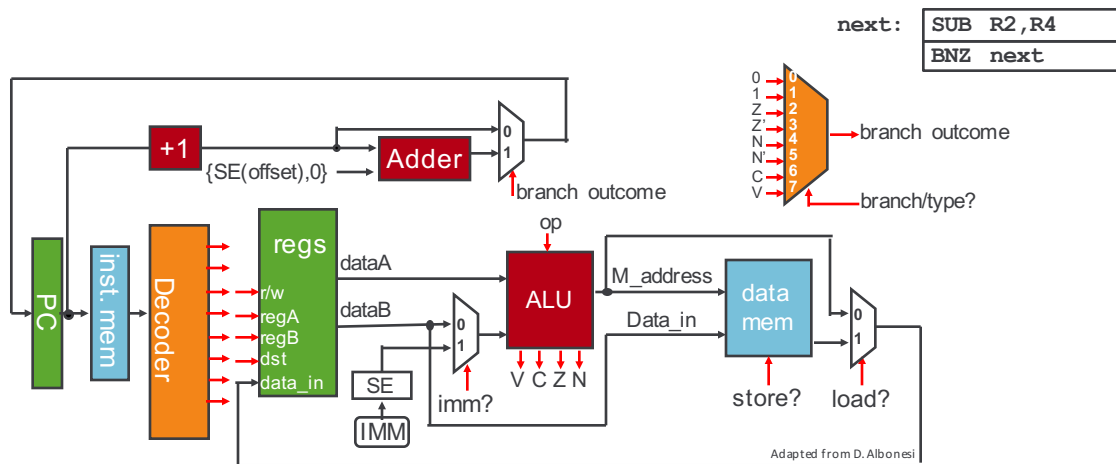
instructions provided by hardware to write a program

Instruction Set Architecture (ISA)

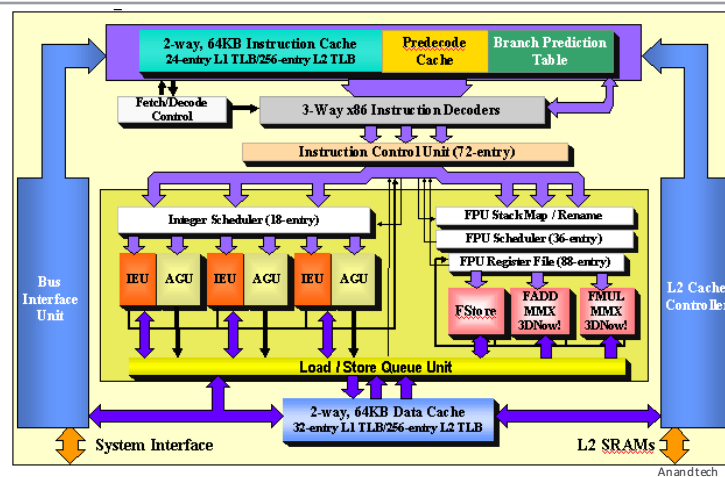
- **Contract between hardware and software**
- Hardware free to implement it in different ways
 - ... as long as software can't tell the difference!
 - Improvements across processor generations
 - Design choices across product families (e.g., high-performance vs. low-energy)
 - High-performance trickery (e.g., out-of-order execution)
- Software free to use any syntax
 - ... as long as it can be translated into working assembly program!



A simple implementation (made up)

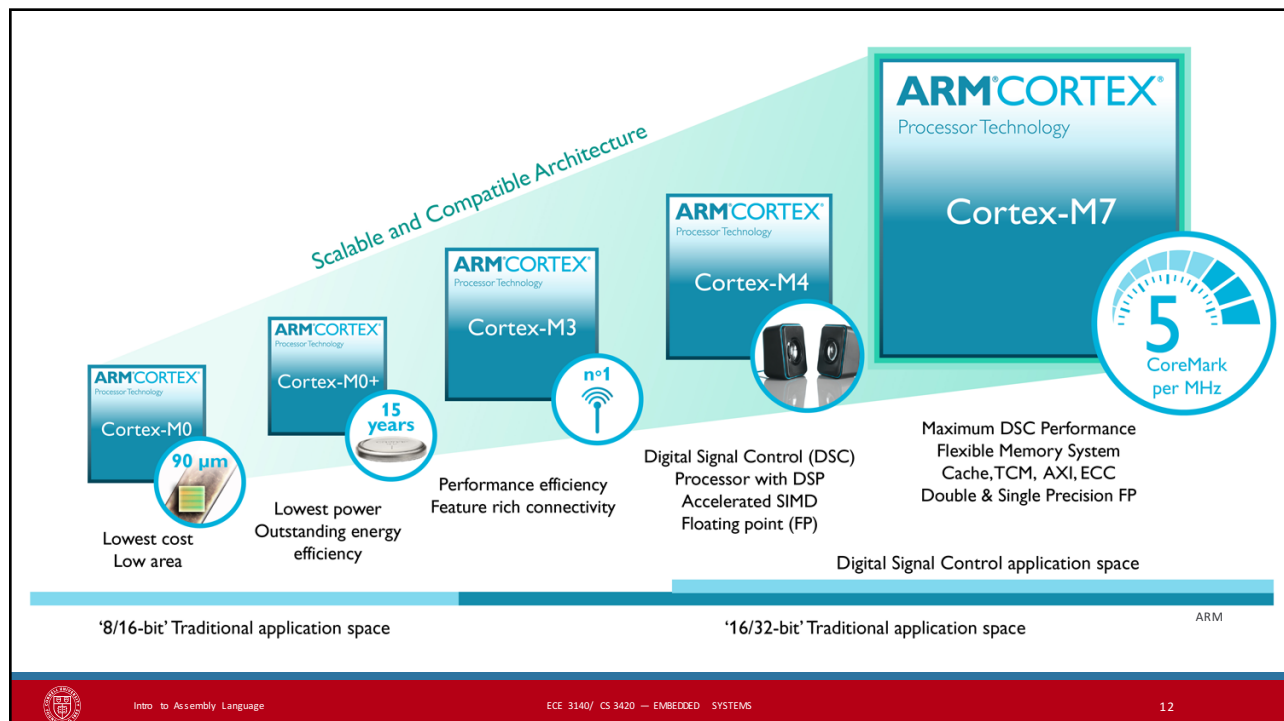


A complex implementation (AMD Athlon)



ARM Cortex-M architecture

- 32-bit datapath (operands and results)
- 32-bit addressing space (memory size)
- *Thumb* ISA (vs. *ARM* ISA in Cortex-Ax)
 - Most instructions 16-bit encoding for compactness
 - Some instructions 32 bits to encode additional functionality
- Different products: M0, M0+, M3, M4, M7
 - “Core” ISA is the same; extensions for functionality
- ARM is *fabless*: License IP, implementation up to customer



Instruction set summary

Instruction Type	Instructions
Move	MOV
Load/Store	LDR, LDRB, LDRH, LDRSH, LDRSB, LDM, STR, STRB, STRH, STM
Add, Subtract, Multiply	ADD, ADDS, ADCS, ADR, SUB, SUBS, SBCS, RSBS, MULS
Compare	CMP, CMN
Logical	ANDS, EORS, ORRS, BICS, MVNS, TST
Shift and Rotate	LSLS, LSRS, ASRS, RORS
Stack	PUSH, POP
Conditional branch	IT, B, BL, B{cond}, BX, BLX
Extend	SXTH, SXTB, UXTH, UXTB
Reverse	REV, REV16, REVSH
Processor State	SVC, CPSID, CPSIE, SETEND, BKPT
No Operation	NOP
Hint	SEV, WFE, WFI, YIELD

ARM



Instruction format

- General format: `op <dst> <src1> <src2>`
 - There may be fewer source operands and/or no destination
 - Operands may be registers, or (sometimes) immediate constants
- Some examples:
 - `SUB R7, R2, R4` (“subtract”)
 - $R7 \leftarrow R2 - R4$
 - `SUBS R2, R2, #3` (“subtract and update status flags”)
 - $R2 \leftarrow R2 - 3$; update status flags (`SUBS` vs. `SUB`) according to result (will cover shortly)



Instruction format

- **General format:** `op <dst> <src1> <src2>`
 - There may be fewer source operands and/or no destination
 - Operands may be registers, or (sometimes) immediate constants
- **Some more examples:**
 - **CMP R2,R4** (“compare”)
 - Update status flags (will cover shortly) according to result of R2-R4; drop result
 - **BNE <label>** (“branch if not equal”)
 - Jump (branch) to instruction at position <label> in the program if status flag Z ≠ 0
 - Operand actually encoded as offset from current position in the program



Operands

- **General-purpose registers**
 - R0-R7 “low registers,” accessible by all instructions
 - R8-12 “high registers,” not accessible by many 16-bit instructions
 - R13-15 reserved for special purposes (will cover shortly)
 - R15 = “program counter” (PC); R14 = “link register” (LR); R13 = “stack pointer” (SP).
 - Write to at your own peril!
- **Immediate values**
 - Encoded within the instruction format
- **Memory locations (will cover shortly)**



Instruction encoding

- 32-bit instruction if bits [15:11] of the first half-word are 0x1d-f
 - Otherwise, 16-bit instruction
- Opcodes must be unambiguous
 - First few bits tell decoder what to expect in the rest of the instruction

15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
opcode															

Table A5-1 shows the allocation of 16-bit instruction encodings.

Table A5-1 16-bit Thumb instruction encoding

opcode	Instruction or instruction class
00xxxx	Shift (immediate), add, subtract, move, and compare on page A5-6
010000	Data processing on page A5-7
010001	Special data instructions and branch and exchange on page A5-8
01001x	Load from Literal Pool, see <i>LDR (literal)</i> on page A6-90
0101xx	Load/store single data item on page A5-9
011xxx	
100xxx	
10100x	Generate PC-relative address, see <i>ADR</i> on page A6-30
10101x	Generate SP-relative address, see <i>ADD (SP plus immediate)</i> on page A6-26
1011xx	Miscellaneous 16-bit instructions on page A5-10
11000x	Store multiple registers, see <i>STM / STMIA / STMEA</i> on page A6-218
11001x	Load multiple registers, see <i>LDM / LDMIA / LDMFD</i> on page A6-84
1101xx	Conditional branch, and supervisor call on page A5-12
11100x	Unconditional Branch, see <i>B</i> on page A6-40

ARM



Example: **sub** (immediate)

Encoding T1 All versions of the Thumb ISA.

SUBS <Rd>, <Rn>, #<imm3>

SUB<c> <Rd>, <Rn>, #<imm3>

15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
0	0	0	1	1	1	1	imm3			Rn			Rd		

d = UInt(Rd); n = UInt(Rn); setflags = !InITBlock(); imm32 = ZeroExtend(imm3, 32);

Encoding T2 All versions of the Thumb ISA.

SUBS <Rdn>, #<imm8>

SUB<c> <Rdn>, #<imm8>

15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
0	0	1	1	1	Rdn			imm8							

d = UInt(Rdn); n = UInt(Rdn); setflags = !InITBlock(); imm32 = ZeroExtend(imm8, 32);

Outside IT block.

Inside IT block.

Outside IT block.

Inside IT block.

Encoding T3 ARMv7-M

SUB(S)<c>, W, <db>, <dn>, #<const>

15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
1	1	1	1	0	1	0	1	0	1	0	1	0	1	0	1
S		Rn		imm3		Rd		imm8							

If Rn == '1111' then SEE OP (immediate);
If Rn == '1101' then SEE SUB (SP minus immediate);
d = UInt(Rd); n = UInt(Rn); setflags = (S == '1'); imm32 = ThumbExpandImm(imm3, imm8);
If d IN {13, 15} || n == 15 then UNPREDICTABLE;

Encoding T4 ARMv7-M

SUBW<c>, <db>, <dn>, #<imm12>

15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
1	1	1	1	0	1	1	0	1	0	1	0	1	0	1	0
S		Rn		imm3		Rd		imm8							

If Rn == '1111' then SEE ADR;
If Rn == '1101' then SEE SUB (SP minus immediate);
d = UInt(Rd); n = UInt(Rn); setflags = FALSE; imm32 = ZeroExtend(imm3, 32);
If d IN {13, 15} then UNPREDICTABLE;

ARM



Big-endian and little-endian are terms that describe the order in which a sequence of bytes are stored in computer memory. Big-endian is an order in which the "big end" (most significant value in the sequence) is stored first (at the lowest storage address). Little-endian is an order in which the "little end" (least significant value in the sequence) is stored first.

most variables are going to be stored in memory


the more memory, the taller it is
Like a line: in a line, indexed/
segmented in a linear way
will store consecutively

Memory organization

- Integer value types: byte (8b), half word (16b), word (32b)
- 32-bit addresses = 4 GB addressing space
 - Addressable by byte
- Words and half words *aligned*
 - E.g., 4-byte word \Rightarrow base address divisible by 4; value stored in locations $\text{base} + \{0,1,2,3\}$
 - Cortex-M typ. *little-endian*: least-significant byte stored in base address (vs. most-significant byte in *big-endian*)
- Example: Write all legally accessible values \longrightarrow

if address is odd, there will be an exception
for short MUST be even
for word MUST be divisible by 4 (4B)

0xaaaaaaaaab	0xef
0xaaaaaaaaa	0xbe
0xaaaaaaaa9	0xad
0xaaaaaaaa8	0xde



Intro to Assembly Language

ECE 3140/ CS 3420 — EMBEDDED SYSTEMS


19

depends on how you access data
the data by itself tells you nothing, the program is what gives meaning
Is number signed/unsigned, string? \rightarrow program dictates choice

Ex. Little endian: word 0xad de
read short if address was at a: ef be
two byte can only read even (can't read 0xad- can't access half words with an odd address)
Can read each individual bytes (ef,be,ad,de)
All Legal:
0xadde, 0xefbe, 0xde, 0xed, 0xbe, 0xef

Load/store operations

- ARM is a load-store architecture
 - Memory values can only be accessed through load/store instructions
 - All data processing takes place in registers
 - Dramatically reduces complexity of ISA and implementation
- LDR <Rt>, <address>: load (32-bit) word in M[address] into Rt
- STR <Rt>, <address>: store Rt's (32-bit) content into M[address]
- Other opcodes for half-word, byte, etc.



Intro to Assembly Language

ECE 3140/ CS 3420 — EMBEDDED SYSTEMS

20

Memory: must load into register and operate
Write: write from register back to memory
Memory: only can use LD and STR
**Simpler ISA: faster, cheaper, BUT need more instructions to carry out same task
BUT total size is still about the same/managable

Loading sub-word data sizes

How to load half-word/byte data into (32-bit) register?

- Unsigned: Pad with zeroes—e.g., 0x82 (130) → 0x00000082
- Signed: Sign extension—e.g., 0x82 (-126) → 0xffff82

	Signed	Unsigned
Byte	LDRSB	LDRB
Half-word	LDRSH	LDRH

ARM

Can also sign-extend sub-word value already in a register:

	Signed	Unsigned
Byte	SXTB	UXTB
Half-word	SXTH	UXTH

ARM



Addressing modes

Addressing modes: Calculate *effective address* on the fly

- Few modes ⇒ simpler ISA and implementation

[R7,0] -> [R7,4] OR [R7, -4] (4B)

■ [<Rn>, <offset>]: effective address is <Rn>+<offset>

- <Rn> is the “base register;” it can be R0-7, PC, or SP
- <offset> can be immediate constant or another register <Rm>

■ [<Rn>, <offset>]!: Write effective address back to base register (“pre-update”)

Use for for loop/ program counter

■ [<Rn>], <offset>: Use base register as effective address, then update base register with newly calculated address (“post-update”)

post update: use Rn as address, calculate next address and put in register



will not program in x,y coordinates
will just use single additions off of base offsets

Condition codes

- Special APSR register holds four one-bit *condition codes*
 - *Application Program Status Register*
 - N: Result of last status-updating instruction was Negative
 - Z: Result of last status-updating instruction was Zero
 - C: Last status-updating instruction produced Carry
 - V: Last status-updating instruction produced oVerflow
- “S” suffix indicates ALU instruction updates APSR
 - E.g., SUB vs. SUBS, ADC vs. ADCS, etc.
 - Compare instructions always update APSR (obviously)
 - compare and compare negative (CMP, CMN)



Pseudo-instructions

- Assembly-like syntax, but emulated
 - Another instruction can accomplish the same
 - Small block of code (less frequent)
 - Part of the ISA specification; sometimes assembler-specific
- Example: `LDR <Rt>, <immediate>`
 - If <immediate> representable with 8 bits, use `MOV <Rt>, <immediate>`
 - Otherwise (one possible solution):
 - Place <immediate> in program's *literal pool* (well-known memory block)
 - Use `LDR <Rt>, [PC, <offset>]` where <offset> indicates position of literal relative to current PC

